

# Traceability: What does it really mean for QVT?

Edward D. Willink<sup>1</sup> and Nicholas Matragkas<sup>2</sup>

<sup>1</sup> Willink Transformations Ltd., UK,  
`ed-at-willink.me.uk`

<sup>2</sup> University of York, UK,  
`nicholas.matragkas-at-york.ac.uk`

**Abstract.** Traceability in Model Transformation languages supports not only post-execution analysis, but also incremental update and co-ordination of repetition. The Query/View/Transformation family of languages specify a form of traceability that unifies high and low level abstraction in declarative and imperative transformation languages. Unfortunately this aspect of the QVT specification is little more than an aspiration. We identify axioms that resolve the conflicting requirements on traceability, and provide a foundation for resolving further issues regarding equality, transformation extension and rule refinement.

**Keywords:** Model Transformation, QVT, Traceability, Transformation Extension, Rule Refinement

## 1 Introduction

The successes of early model transformation prompted the Object Management Group (OMG) to issue a Request for Proposal for a standard solution to the problems of querying, viewing and transforming models. The initial submissions eventually converged to give us the Query/View/Transformation (QVT) family of languages. The specification standardizes research work and so the specification is imperfect.

Problems in an OMG specification are addressed by a Revision Task Force (RTF). The QVT 1.2 RTF successfully resolved most of the outstanding problems but a small number were deferred for lack of time, a convenient euphemism for help-needed. The vagueness of the traceability aspects of the specification were addressed in an earlier version of this paper. Following some discussion, this paper provides a more detailed consideration of what traceability needs to achieve and the approach take for QVT 1.3.

In Section 2 we first summarize the specified aspects of QVT traceability, then in Section 4 we dig down to the underlying purpose of traceability in a declarative context and formulate some axioms that enable us to consider the Object-Orientedness of transformations. The more challenging context of Imperative Transformations is discussed in Section 5. In Section 6 we outline a Trace Model that can be used more realistically throughout QVT. We present related work in Section 7 and finally conclude in Section 8.

## 2 QVT Traceability as Specified

The QVT specification defines a family of three languages, in part reflecting differing viewpoints amongst the original submitters and in part reflecting different community requirements.

- QVT Operational Mappings (QVTo) an imperative transformation language.
- QVT Relations (QVTr) a high level declarative transformation language.
- QVT Core (QVTc) a low level declarative transformation language.

These languages are unified via the QVTc semantics. QVTr transforms directly to QVTc. QVTo and arbitrary externally supplied blackboxes are defined with respect to an equivalent QVTr Relation. The practical compatibility of this unification has yet to be demonstrated, since the various QVTo and QVTr vendors find sufficient challenges in realizing a primary language and blackboxes, without worrying about the alternative paradigm. There are no complete QVTc implementations.

In this paper we concentrate on a narrower and potentially much clearer form of compatibility; the *trace records* which are supposedly common to all three languages.

One benefit and consequence of the low level QVTc language is that the user can and must specify the maintenance of the classes that establish traceability. These form the middle model for a typical bidirectional QVTc rule in which left and right hand models are related with the aid of the intermediate trace model.

A `Package2Schema` mapping may therefore operate between a `Package` *p* and a `Schema` *s* while maintaining a *trace record* of the intermediate `TPackage2Schema` *trace class* with *p* and *s* properties.

```
class TPackage2Schema {
    property p : Package;
    property s : Schema;
}
```

When the mapping is executed in the forward direction, the *p* property is initialized with the matched input `Package` and the *s* property with the generated output `Schema`. The *trace instance* records the `Package2Schema` relationship between *p* and *s* when the source-to-target relationship is first established. Thereafter, another mapping can locate the schema corresponding to `aPackage` by the OCL navigation `aPackage.TPackage2Schema.s`. This starts at `aPackage` and uses the `Package::TPackage2Schema` relationship, which is the implicit opposite of the `TPackage2Schema::p` relationship. This locates the *trace record* and then the `TPackage2Schema::s` relationship navigates to the target. This is a demonstration of the power of an explicitly modeled *trace class*; there is no need for any special operators such as `resolve` in QVTo or `resolveTemp` in ATL.

The need for manual maintenance of the *trace records* is eliminated in QVTr for which a mapping to QVTc is suggested by Rule 1 of Clause 10.2 of the

QVT specification. Unfortunately this rule neglects to specify an algorithm for construction of the trace class name and so there is no guidance on how name clashes are to be avoided.

The lack of any common base class to support polymorphic *trace instance* maintenance of *trace classes* makes this aspect of the QVT specification totally unsuitable for practical tooling. In Section 6 we outline a more rational proposal.

The specification of QVTo *trace classes* is distinctly vague since it relies on the underspecified correspondence between a QVTo MappingOperation with an unspecified Operation signature, and a QVTr Relation. Overloading is even more of a problem in QVTo with its disjunct mappings. The need to extend the *trace data* to accommodate *in/inout/out* parameters is identified without any clarification on how the trace data is modified in a way that remains compatible with QVTc. *inout* parameters present a further challenge; in QVTo an *inout* parameter has distinct *in* and *out* values. This is incompatible with a one-to-one correspondence with a QVTr relation. QVTr is multi-directional, so a parameter may be *in* or *out* depending on the direction; it cannot be both *in* and *out*.

### 3 Traceability Use Cases

Traceability is too often referred to as a good thing without ever defining what it is or what it is for, so we will first examine some alternative use cases.

#### 3.1 Requirements Traceability

Anyone from the Systems Engineering community is easily confused by the Modeling Transformation usage, since traceability is Requirements Traceability; the correlation of downstream deliverables with upstream requirements. Model transformations can fulfil system requirements either as tooling or as deliverables and so there is a double opportunity for confusion. In practice Requirements Traceability can be important, but it is nothing to with the traceability concerns of model transformations.

#### 3.2 printf Traceability - What happened?

The simplest form of model transformation traceability answer the question “What happened?”. It provides voluminous diagnostic output to facilitate debugging or performance analysis of transformation execution. This may just require numerous printf statements, but since we are modeling we may at least hope that the voluminous output takes the form of a model.

#### 3.3 Debugger Traceability - What is happening?

Use of a debugging tool helps answer the question “What is happening?”. This now requires some interaction between the execution state and typical source level stepping and variable browsing capabilities.

### 3.4 Internal Traceability - Has X executed?

Model transformation languages generally prohibit re-execution of the same rule. We use rule here; other languages may call the basic model transformation language building block a mapping or a relation. Prohibiting re-execution requires a recording of the prior execution so that it can be detected and re-used.




**Fig. 1.** Simple Rule Execution

Fig 1 shows a simple Rule named R1, whose invocation with an input instance s1 of type S1 produces an output instance t1 of type T2. If R1 is invoked again with s1 as input, we want to inhibit the re-execution and use the previous t1 output.

### 3.5 Re-use Traceability - What was the result of the X execution?

Re-using an execution is perhaps just a problem for the implementors of a transformation tool. However a slightly different form of re-use is critical to useable model transformation.



**Fig. 2.** Rule Execution with Re-use

Fig 2 shows another Rule named R2 that converts a parent S1 object and a child S2 object into a parent T1 object and a child T2 object. However, we want to re-use the t1 result previously produced by the R1 rule for the s1 instance, rather than create another t1.

The ability to interrogate a rule execution is so important that modeling languages provide a special language construct. *resolveTemp* in ATL, *equivalent* in Epsilon, *resolve* in QVTo, *when* in QVTr.

### 3.6 Incremental Traceability - What do I need to re-execute?

The most powerful use of traceability is to support proportionate re-execution of a transformation after changes to input models.

In the absence of white box knowledge, re-execution, of the Fig 2 black box, is necessary if either *s1* or *s2* has changed. If this re-execution is able to re-use the in-memory models, orchestrating the re-execution is relatively straightforward; the *s1* and *s2* objects can observe the changes and schedule re-execution. However when the re-execution occurs at a different time or place to the original, a detailed recording of the original execution must be persisted for use by the re-execution. The recording, *trace representation* of *s1* and *s2* must be compared with the prevailing *s1* and *s2* to determine whether change has occurred.

## 4 Declarative QVT Traceability Purpose

The QVT specification aspires to support incremental re-execution, so we examine what properties we need traceability to exhibit in order to perform declarative transformations. Once we understand declarative transformations, we can then see in Section 5, how the QVTo imperative transformation language can conform to the same principles.

### 4.1 Trace Record Concepts

**Trace Representation** A *trace representation* summarizes the state of an input or an output model element.

The *trace representation* of any pair of model elements contains sufficient information to determine whether the two model elements are equal. Additionally an *output trace representation* contains sufficient content to allow the represented element to be reconstructed.

UML, and consequently OCL and transformation languages based on UML or OCL, define two importantly different forms of model element.

Each Class instance or object represents a distinct model element; no two objects are ever equal. It is therefore convenient to summarize the state of any object by its memory address since each object has a distinct address. Reconstruction of the object from the summary is trivial, so long as the object remains in memory. Reconstruction from a persisted form requires a corresponding *trace representation* of each object property.

An instance of a *DataType* (such as *Integer* or a *Collection(String)*) has a value. Instances of *DataTypes* are equal when they have the same value, even if the instances are stored at distinct addresses. 4 is clearly always equal to 4 no matter what address the 4 is at. For values such as *Integers*, whose storage requirements are small, it is convenient to use the entire value as the representation. For potentially very large values such as *Strings* or *Collections*, it is convenient to use the address as the summary, making reconstruction easy, but requiring care to compare the underlying value rather than just the *trace representation* when determining whether two *DataType* values are equal.

**Trace Record** A *trace record*, which may also be called a *trace instance* of a *trace class*, comprises

- the identity of the rule that is traced
- an ordered list of the *input trace representations* of each rule input argument
- an ordered list of the *output trace representations* of each rule output result

(The ordering of the lists may be established by some deterministic algorithm such as an alphabetic sort applied to the names of the rule parameters.)

The trace record for Fig 2 may be written as the tuple  $\{R2, \{s1, s2\}, \{t1, t2\}\}$ .

**Trace Identity** A *trace identity* comprises the input elements of the trace record; rule identity and *input trace representations*. The *trace identity* of Fig 2 is  $\{R2, \{s1, s2\}\}$ . Traceability may also be considered in the inverse direction; the *inverse trace identity* is  $\{R2, \{t1, t2\}\}$ .

## 4.2 Trace Record Characteristics

The preceding traceability concepts exhibit the following characteristics to provide useful traceability.

**Trace Data** The *trace data* comprises all the *trace records*.

**Trace Identity Uniqueness** No two *trace records* share the same *trace identity*.

**Trace Record Creation** Every rule execution has a *trace identity* and a corresponding *trace record*.

Consequently any attempt to execute a rule that would require a duplicate *trace identity* must be detected and suppressed. The results of the suppressed execution are reconstructed from the *output trace representations* of the previous execution.

**Trace Record Access** The *trace data* may be exploited to recover a *trace record* for a given *trace identity* and so locate the target for a graph-to-graph transformation.

Consequently the *output trace representation* must be consistently determined by the *trace identity*.

If an underlying input model element changes, then the *input trace representation* changes too and so the rule must be re-executed if the corresponding outputs are required.

### 4.3 Tracing immutable Class instances and DataType values

With distinct input and output models, the input model elements of a declarative transformation are stable and so the use of an address to summarize each unchanging object is valid. This incurs very limited cost apart from limiting Garbage Collection in long running programs. This may also be possible for well-behaved imperative transformations that use only one state of each object.

### 4.4 Tracing mutable Class instances and DataType values

However for an in-place transformation or less well-behaved imperative transformation, model elements may be updated requiring a more complete *input trace representation* to ensure that the rule is executed again for a changed input to yield a correspondingly changed output.

**In-place declarative transformations** An in-place transformation is equivalent to an out-of-place transformation followed by a copy of the transformed results back to the input. This avoids the need to trace multiple objects states. In practice an efficient implementation may sequence object changes to allow direct modification of the input objects. The optimization of object modifications just needs to be mirrored by the trace record maintenance.

**Imperative transformations** Maintaining multiple DataType values for an imperative transformation is comparatively straightforward. It just requires sufficient memory to accommodate the numerous similar deep clones, or some more efficient but complex mechanism to share memory between similar values.

Maintaining multiple states of an Object is very much harder. Naively, a shallow object clone might seem more than sufficient, however MOF imposes a limitation that no object can have more than one container. Consequently, the clone cannot re-use containment relationships with other objects without corrupting the original object; the child-stealing phenomenon [11]. A deep clone solves the child-stealing problem at the expense of a large amount of memory per object state and a new difficulty comparing cloned objects for equality. In practice some form of versioning must be used to provide efficient representations of the various states without the impediments of MOF objects.

### 4.5 Transformation Extension and Rule Refinement

All three QVT languages define transformations that may extend other transformations, and rules that may refine other rules, but neglect to define what this may mean. The immediate question is are transformations Object-Oriented and do they observe the Liskov Substitution Principle[7]?

We would very much like transformations to be Object-Oriented so that we can use familiar derive and override practices, so let us assume a strong analogy between the new Transformation/Extension/Rule/Refinement concepts

and the familiar Class/Inheritance/Operation/Overriding. We will exploit the traceability axioms to see what limitations we may need.

The Liskov Substitution Principle is of course never observed unequivocally, but disciplined transformation writers can define those behaviors for which substitutability is important and code their refinements accordingly.

**Transformation Extension** For the model parameters of an extended transformation, we might look for an analogy with the arguments of a derived constructor, where the arguments can be very different and must be converted to use one of the inherited constructors. However for transformations there is just a single interface of one model-type per model-parameter and there are no conversion facilities. A model-type is a list of packages of types that the transformation compiler needs to consider; the model-type is not significant at run-time. The transformation extension therefore just routes the models of the extending transformation to the extended transformation. Since the execution direction is unknown, every type that could be generated by the extended transformation must be included by the model-type of the extending transformation, so we require:

*All extending transformation model-types must be super-sets of the extended model-types.*

**Rule Refinement** The declarative QVT languages are multi-directional and so we do not know which rule parameters are inputs and outputs. There is therefore no opportunity for covariant or contra-variant refined rule parameter types. However rules have a very important execution distinction from operations. When an operation is invoked, the operation or its overload executes. However when a rule is invoked rule execution is conditional on the satisfaction of its predicates; the parameter types are therefore just a stylized predicate. Wimmer et al. [12] observe that for declarative rules a refining rule may even have additional parameters, or may refine multiple rules. This is a consequence of the rule invocation occurring as a consequence of a pattern match. In contrast, an imperative rule is explicitly invoked and so its refinements must be argument count compatible with the explicit invocation.

**Dynamic Dispatch** For OO-style dynamic dispatch of rules to work, we require:

*QVTc and QVTr must have a this object that is an instance of the executing transformation.*

The instance can be stateless.

**Tracing** When we consider a *trace record* involving execution of a refined rule, should the apparent invoked rule or the actual executed rule be recorded? A subsequent user of the trace must be able to resolve against the apparent rule,



rather than all possible refinements, but must also be able to resolve against the actual rule where that is known.

*There must be a trace record for each possible trace identity.*

This does not necessarily imply multiple *trace records* and in Section 6 we propose to allow a *trace record* to have multiple *trace identities*.

**Composition** An overridden declarative rule provides a composition of predicates and assignments.

In accordance with Design by Contract[8] all predicates should be satisfied. There may obviously be scope for optimization of redundancy.

Similarly all assignments should be made, but since multiple assignments may conflict, each most derived assignment should be chosen, thereby allowing a refining rule to adjust as well as extend the behaviour of a refined rule.

*A rule refinement hierarchy is equivalent to a single composite rule.*

**Concurrency** If we have a system in which more than one copy of a transformation executes on the same input models, we may want to share common model elements but must keep composed model elements exclusive. Providing a single overall *trace data* will satisfactorily share the common model elements, but we must avoid Child Stealing of the exclusive model elements. Most rules create composed children with the intended parent as one of the traced inputs, so provided the parents are exclusive, the children will be too. We therefore only need to ensure that the root parent is exclusive, which may be achieved by using the transformation *this* as part of the *trace identity*.

*Trace data may be shared by concurrent transformations.*

This same reasoning applies to the nested execution of extended/extending transformations or refined/refining rules.

## 5 Imperative Traceability

So far we have concentrated on the declarative QVT languages where the ability for objects to mutate and the ability to impose a schedule is very limited. Declarative scheduling is an implicit consequence of the inability to perform a rule until its inputs are available. Complex input/output dependencies may result in the rules being grouped into passes. Object mutation is limited to one value at the input/output of each ‘pass’.

### 5.1 Imperative Characteristics

In contrast QVTo allows the user to

- specify an explicit execution order
- make arbitrary changes to objects
- use mutable collections

- exploit changeable global context
- exploit changeable transformation context
- have *inout* mapping arguments

Of these only the explicit order is not a problem for traceability; it just allows an infeasible order to be programmed.

**Object changes** If an object changes, the result obtained from a rule execution using that object prior to the change may not be valid and so the assumption that its state can be summarized by just its address is unsound; the *input trace representation* must be extended to all mutable fields that influence the rule execution. This may be difficult to determine when complex helper functions are used, so it may be necessary for the *input trace representation* to create a deep-clone of the traced input unless an existing deep-clone can be re-used.

The *trace record* must also maintain a deep-clone of the output in order to avoid any corruption by user assignments, but this makes it impossible to return the same result as a previous mapping execution, since each new result must be a clone to avoid corruption by assignments.

**Collection changes** QVTo introduces two new forms of mutable collection: List and Dict. These are very convenient for imperative programming, but have never been fully characterized in terms of their UML alignment. Is List{1} equal to List{1}? Obviously Yes. Is List{2} equal to List{1}? Obviously Not, not even if List{2} arose from changing List{1}. List is therefore clearly a Mutable DataType whose memory location is irrelevant since only its value is interesting.

This therefore presents the same problem as for Object changes. The *input trace representation* and *output trace representation* must be deep-clones to avoid value corruption in the *trace record*.

**Global context changes** The global context such as configuration properties may influence the mapping execution and the global context may be changed between mapping executions. It is therefore necessary to include the prevailing state of all relevant mutable global context as part of the *trace identity* in order to accurately determine whether the previous mapping execution result can be re-used.

**Transformation context changes** Transformation context such as contextual properties and intermediate classes may also influence the mapping execution and again this context may be changed between or during executions. It is therefore also necessary to include the prevailing state of all relevant mutable transformation context as part of the *trace identity*.

**Mapping argument changes** A QVTo mapping may have *inout* arguments which may influence the behavior, so we should deep-clone the relevant parts

of the value on input as one of the *input trace representations* and the whole output as one of the *output trace representations*.

## 5.2 Assessment

It may be noted that all this additional context is also required to fulfill the goal that the trace record contains the information necessary to determine what needs to be re-executed as part of an incremental update.

QVTo is intended to be a practical language, so introducing all the additional *trace record* content identified above to make traceability sound seems unacceptable. The cost of all the deep-cloning is obviously large and may be quadratically so. Consider a relatively simple mapping in which the programmer either does not understand, or does not trust, the built-in *trace record* resolution capabilities. An *inout* Dict is passed in order to keep track of all the input to output mappings. This Dict will grow in size with each mapping and so we have a steadily growing Dict to deep-clone and consequently  $O(N*N)$  complexity.

We must instead try to tighten to the language to retain utility while adding some integrity.

The major problems of deep-cloning to stabilize the *trace identity* can be avoided if all traced inputs are transitively immutable. This can be achieved for inputs by requiring an *in* rather than *inout* parameter and introducing a new restriction that only constants, exclusive clones or *in* parameters may be passed to *in* parameters. This guarantees that *in* can be used without deep-cloning.

A similar policy could be applied to outputs, but the language would be useless; it would be impossible to execute a mapping to create a parent object and then populate its children, either in the caller or in a separate piece of code that uses the trace to resolve the output for update in a second pass. It seems necessary to allow the *output trace representation* to evolve.

Changes to the global context can be ignored, in the hope that users will restrict the use of global context to stable configuration that requires a total re-evaluation for any change. It may be possible for tooling to detect global context sensitivity and guide users into refactoring their transformation so that unstable global context becomes a disciplined input model.

Some changes to the transformation context could be similarly ignored, but the transformation itself cannot. There is no way that two concurrent imperative transformations operating on the same models can safely share mapping execution outputs. The transformation *this* must therefore form part of an overall *trace identity* or distinct *trace data* must be maintained for each transformation execution.

*inout* arguments must also be ignored in the hope that their usage is to accumulate additional information that does not influence the mapping execution result. This is true of the user-maintained traceability example above. It may often be possible to verify this statically, but not always. We may perhaps be able to impose a transitive prohibition on an *in* or *out* argument ever being used as an *inout* argument.

### 5.3 Transformation Extension and Mapping Refinement

The declarations of QVTo’s `OperationalTransformation` and `MappingOperation` have many similarities to the declarative `Transformation` and `Mapping` so we may look to achieve Object Oriented behavior for imperative transformations as well.

QVTo provides additional mapping refinement options.

A *disjunct* mapping provides an outer mapping which redirects to one of a number of inner mappings according to a type and guard-based selection. Since the outer mapping performs no object creation, the *trace record* for a disjunct mapping can be provided by the selected inner mapping, augmented by the additional *trace identity* of the also-executed outer mapping.

An *inherited* mapping creates an object in the outer mapping before the inherited inner mapping initializes it and returns control for further execution by the outer mapping. Since the outer mapping performs the object creation, the *trace record* for an inherited mapping can be provided by the outer mapping, augmented by the additional *trace identity* of the also-executed inner mapping.

A *merged* mapping executes an outer mapping then executes an inner mapping. Since the outer mapping performs the object creation, the *trace record* for an inherited mapping can be provided by the outer mapping, augmented by the additional *trace identity* of the also-executed inner mapping.

### 5.4 Incremental Update

The QVT specification mentions the utility of the *trace data* to support efficient incremental update. This may be supported by QVTc and consequently QVTr, but is certainly not for QVTo as specified.

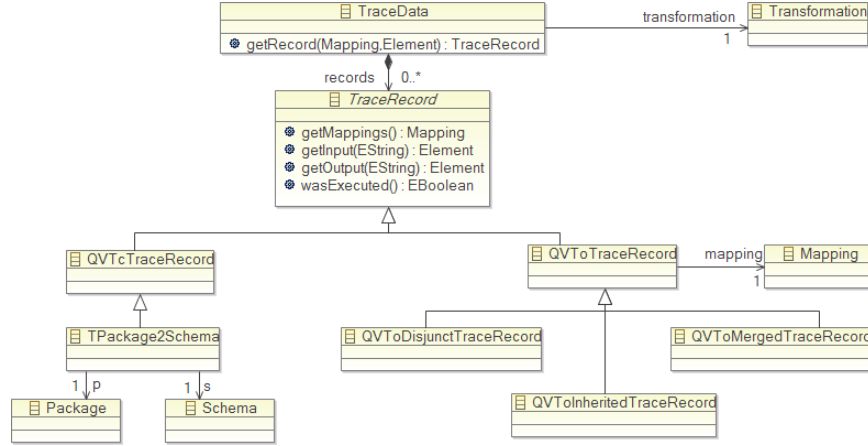
QVTo specifies that a *trace record* is created during the object initialization phase, which occurs after guards have been evaluated. Consequently there is no trace of model elements that were not created and so it is difficult for an incremental update to correctly handle a change that affects the execution of the guards.

QVTo provides no *trace record* for explicit object creation or for explicit cloning.

These limitations in conjunction with the pragmatic exclusion of global, transformation and *inout* context from the *trace record* suggest that in order to achieve reliable incremental update with QVTo, it may be necessary to impose so many declarative characteristics that it may be more profitable to develop a QVTo to QVTr migration assistant than an incremental QVTo solution.

## 6 Multi-language solution

In Section 2 we identified the strong bias of the QVT specification trace model to QVTc. In Figure 3 we suggest a more flexible metamodel that avoids imposing enumerated QVTc-style trace classes on QVTo, supports polymorphic access to



**Fig. 3.** Metamodel for the Trace Data solution

the trace and provides for extension to support specialized tracing of specialized mappings such as QVTc disjuncts.

The abstract **TraceRecord** provides the freedom for **QVTcTraceRecord** and **QVToTraceRecord** to pursue different representation strategies.

The **TPackage2Schema** is an example of an enumerated QVTc trace class. Apart from the inheritance, it is exactly what, and all that, the QVT specification defines.

The **TraceData::getMapping** API supports discovery of a trace record for a given **Mapping** operation and list of **Elements**. The **TraceRecord::getMappings** provides the list of all **Mapping** operations executed. **TraceRecord::getInput** or **getOutput** enable a particular input or output model element to be recovered. **TraceRecord::wasExecuted** enables additional tracing of unexecuted mappings to be recorded to support incremental update.

## 7 Related Work

In the context of model transformation, traceability can be used in a variety of scenarios such as impact analysis or object resolution. Due to this the majority of model transformation languages provide direct or indirect support for tracing.

The Atlas Transformation Language (ATL) [5] uses tracing to resolve the interactions and dependencies between the rules of a transformation. The traceability mechanism offered by ATL is implicit, in the sense that the tracing information is captured automatically by the transformation engine without any input or guidance from the end user. Every time a transformation rule is matched, a new trace link object is created between the source element and its corresponding target element(s) by using the native type *ASMTransientLink*. This trace link is assigned the name of the rule, the source and target elements, and it is added to

a link collection named *ASMTransientLinkSet*. This link collection is used internally by the ATL virtual machine and it is not directly accessible by the end user. ATL provides the *resolveTemp()* method, which makes possible to point from an ATL transformation rule to any of the target model elements by using the name of a model element. One limitation of this approach is that the end user does not have richer access to the traceability information. For example, searching by type or iterating over all trace links is not supported. [13] propose a method that allows richer run-time access to traceability information by extending the *ASMTransientLink* and *ASMTransientLinkSet* classes. A second limitation of the ATL traceability support is the fact that after the execution of the transformation the traceability information is discarded. In [4] the authors propose a method for persisting traceability information. In this method traceability-specific code is embedded in the transformation code. When a transformation is executed this code generates a traceability model. The traceability-specific code can be added either manually or by the use of a Higher Order Transformation (HOT). Although the proposed solution solves the problem of persistence of traceability information, it creates an overhead problem.

Another transformation language, which provides direct support for traceability, is the Epsilon Transformation Language (ETL) [6]. Similarly to ATL, ETL uses traceability for resolution of source elements in the target models during a transformation execution. ETL traceability is implicit. When a rule is matched, the ETL engine generates an instance of a *Transformation* class. This instance captures the source element of the rule, a collection of the target elements, and the rule, which was used to create this trace. Once a trace is created it is added to the *TransformationTrace* object, which holds a collection of Traces for a particular transformation. ETL provides the *equivalent()* built-in operation, which uses the generated trace links to automatically resolve source elements to their transformed counterparts. When the *equivalent()* operation is applied on a single source element, it inspects the established transformation trace and invokes the applicable rules (if necessary) to calculate the counterparts of the element in the target model. Epsilon defines a set of model management Ant tasks for orchestration workflows. One of the provided tasks is the *ETL-Task*, whose *exportTransformationTrace* attribute enables developers to export an internal transformation trace to the project context. Finally, ETL supports rule inheritance. When a rule extends another rule, the ETL engine keeps only the sub-rule as part of the trace, ignoring the super-rule.

In Kermeta, a traceability framework for facilitating the trace of model transformations is defined [3]. The framework is built atop a language independent traceability metamodel. In Kermeta, a model transformation trace is defined as a bipartite graph with source and target nodes. To support transformation chains, Kermeta defines every trace as an ordered set of trace steps, each of them representing a single transformation. A trace step can consist of many links, which relate source and target objects. To generate trace information during the execution of a Kermeta transformation, traceability-specific code has to

be embedded in the transformation code. Finally, the generated trace links can be serialized and subsequently used in further model management tasks.

In the Simple Transformer (SiTra) tool, traceability is inspired by the traceability support of QVT. Tracing consists of the *ITrace* interface, which holds a collection of *TraceInstances*. Each trace represents a mapping between a source and a target model element through a transformation rule. An implementation of the *ITrace* interface provides a number of methods to query the trace collection and return all target instances. The Sitra transformation engine ensures that, for each transformation rule being executed, a trace is recorded.

In the context of QVT, [2] argue that the trace mechanism of the QVT Operational does not capture enough information for common scenarios requiring traceability as input. To enrich the information content of this trace, they propose an alternative metamodel which can be used in conjunction with QVT and it can support richer traceability models.

In the approaches presented so far, there is dedicated support for the creation and usage of traceability. In addition to the aforementioned model transformation languages, there are other languages which do not support directly traceability such as AGG [9], VIATRA [10] and GReAT [1]. Such approaches can support traceability indirectly by creating and manipulating trace links as any other element.

## 8 Conclusion

We have identified the intent that a single form of traceability should be shared across the family of QVT languages in order to avoid repeated mapping execution and found this to be incompatible with the current specification.

We formulated the intent of traceability as axioms and identified a potential conflict with composition relationships leading to the Child Stealing phenomenon. We have used the axioms to clarify the semantics of transformation extension, mapping refinement, tracing and concurrency in a declarative context.

We have identified significant limitations in the ability to trace reliably in an imperative context and proposed limitations that QVTo could impose to improve tracing integrity.

**Acknowledgments** Thanks to Adolfo Sanchez-Barbudo Herrera for some helpful comments on an early draft.

## References

1. A. Agrawal, G. Karsai, and F. Shi. Graph transformations on domain-specific models. *Journal on Software and Systems Modeling*, 2003.
2. V. Aranega, A. Etien, and J.-L. Dekeyser. Using an alternative trace for QVT. *Electronic Communications of the EASST*, 42, 2011.

3. J.-R. Falleri, M. Huchard, C. Nebut, et al. Towards a traceability framework for model transformations in kermeta. In *ECMDA-TW'06: ECMDA Traceability Workshop*, pages 31–40, 2006.
4. F. Jouault. Loosely coupled traceability for ATL. In *Proceedings of the European Conference on Model Driven Architecture (ECMDA) workshop on traceability, Nuremberg, Germany*, volume 91. Citeseer, 2005.
5. F. Jouault and I. Kurtev. Transforming models with ATL. In *Satellite Events at the MoDELS 2005 Conference*, pages 128–138. Springer, 2006.
6. D. S. Kolovos, R. F. Paige, and F. A. Polack. The Epsilon transformation language. In *Theory and Practice of Model Transformations*, pages 46–60. Springer, 2008.
7. B. Liskov. Keynote address - data abstraction and hierarchy. In *ACM SIGPLAN Notices 23 (5)*, pages 17–34. ACM, 1988.
8. B. Meyer. Design by Contract. *Technical Report TR-EI-12/CO, Interactive Software Engineering Inc.*, 1986.
9. G. Taentzer. AGG: A graph transformation environment for modeling and validation of software. In *Applications of Graph Transformations with Industrial Relevance*, pages 446–453. Springer, 2004.
10. D. Varró. *VIATRA: visual automated model transformation*. PhD thesis, Thesis, Department of Measurement and Information Systems, University of Technology and Economics, Budapest, 2003.
11. E. Willink. QVT Traceability : What does it really mean? <https://www.eclipse.org/mmt/qvt/docs/ICMT2014/QVTtraceability.pdf>.
12. M. Wimmer, G. Kappel, A. Kusel, W. Retschitzegger, J. Schönböck, W. Schwinger, D. Kolovos, R. Paige, M. Lauder, A. Schürr, and D. Wagelaar. Surveying Rule Inheritance in Model-to-Model Transformation Languages. *Journal of Object Technology*, 11:1–46, 2012.
13. A. Yie and D. Wagelaar. Advanced traceability for ATL. In *1st International Workshop on Model Transformation with ATL*, pages 78–87, 2009.